

miniproject

October 2, 2024

```
[164]: # This Python 3 environment comes with many helpful analytics libraries
        ↳ installed
        # It is defined by the kaggle/python Docker image: https://github.com/kaggle/
        ↳ docker-python
        # For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list
↳ all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that
↳ gets preserved as output when you create a version using "Save & Run All"
# You can also write temporary files to /kaggle/temp/, but they won't be saved
↳ outside of the current session
```

```
[165]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
warnings.filterwarnings("ignore", category=UserWarning)
df = sns.load_dataset('titanic')
```

```
[166]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
#   :-----:-----:-----:
#   Survived        891 non-null    bool
#   Sex             891 non-null    object
#   Age            891 non-null    float64
#   SibSp           891 non-null    int64
#   Parch          891 non-null    int64
#   Embarked        891 non-null    object
#   Fare           891 non-null    float64
#   Cabin          294 non-null    object
#   Ticket         891 non-null    object
#   PassengerId     891 non-null    int64
#   Name           891 non-null    object
#   Sex            891 non-null    object
#   Age           891 non-null    float64
#   Survived      891 non-null    bool
```

```

---  -----  -----  -----
0  survived      891 non-null  int64
1  pclass        891 non-null  int64
2  sex           891 non-null  object
3  age           714 non-null  float64
4  sibsp         891 non-null  int64
5  parch         891 non-null  int64
6  fare          891 non-null  float64
7  embarked      889 non-null  object
8  class         891 non-null  category
9  who           891 non-null  object
10 adult_male    891 non-null  bool
11 deck         203 non-null  category
12 embark_town  889 non-null  object
13 alive        891 non-null  object
14 alone        891 non-null  bool
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB

```

```
[167]: print(df.head())
```

```

   survived  pclass    sex  age  sibsp  parch   fare embarked  class \
0         0      3   male  22.0     1     0   7.2500          S  Third
1         1      1  female  38.0     1     0  71.2833          C  First
2         1      3  female  26.0     0     0   7.9250          S  Third
3         1      1  female  35.0     1     0  53.1000          S  First
4         0      3   male  35.0     0     0   8.0500          S  Third

   who  adult_male  deck  embark_town  alive  alone
0  man         True  NaN  Southampton    no  False
1 woman        False    C    Cherbourg  yes  False
2 woman        False  NaN  Southampton  yes   True
3 woman        False    C  Southampton  yes  False
4  man         True  NaN  Southampton    no   True

```

```
[168]: # Check for missing values
print(df.isnull().sum())
```

```

survived      0
pclass        0
sex           0
age          177
sibsp         0
parch         0
fare          0
embarked      2
class         0
who           0

```

```

adult_male      0
deck            688
embark_town      2
alive           0
alone           0
dtype: int64

```

```

[169]: # Summary statistics
print(df.describe())

```


	survived	pclass	age	sibsp	parch	fare
count	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```


[170]: import seaborn as sns
import matplotlib.pyplot as plt

# Count plot for 'Survived'
sns.countplot(x='survived', data=df)
plt.title('Survival Count')
plt.show()


```




```
[171]: # Count plot for 'Pclass'
sns.countplot(x='pclass', data=df)
plt.title('Passenger Class Distribution')
plt.show()
```




```
[172]: # Survival rate by class
sns.barplot(x='pclass', y='survived', data=df)
plt.title('Survival Rate by Passenger Class')
plt.show()
```




```
[173]: # Count plot for 'Sex'
sns.countplot(x='sex', data=df)
plt.title('Gender Distribution')
plt.show()
```




```
[174]: # Survival rate by gender
sns.barplot(x='sex', y='survived', data=df)
plt.title('Survival Rate by Gender')
plt.show()
```




```
[175]: # Histogram for 'Age'  
df['age'].plot(kind='hist', bins=30, color='lightblue')  
plt.title('Age Distribution')  
plt.show()
```


```
[176]: # Box plot for 'Age'
sns.boxplot(x='survived', y='age', data=df)
plt.title('Age vs Survival')
plt.show()
```



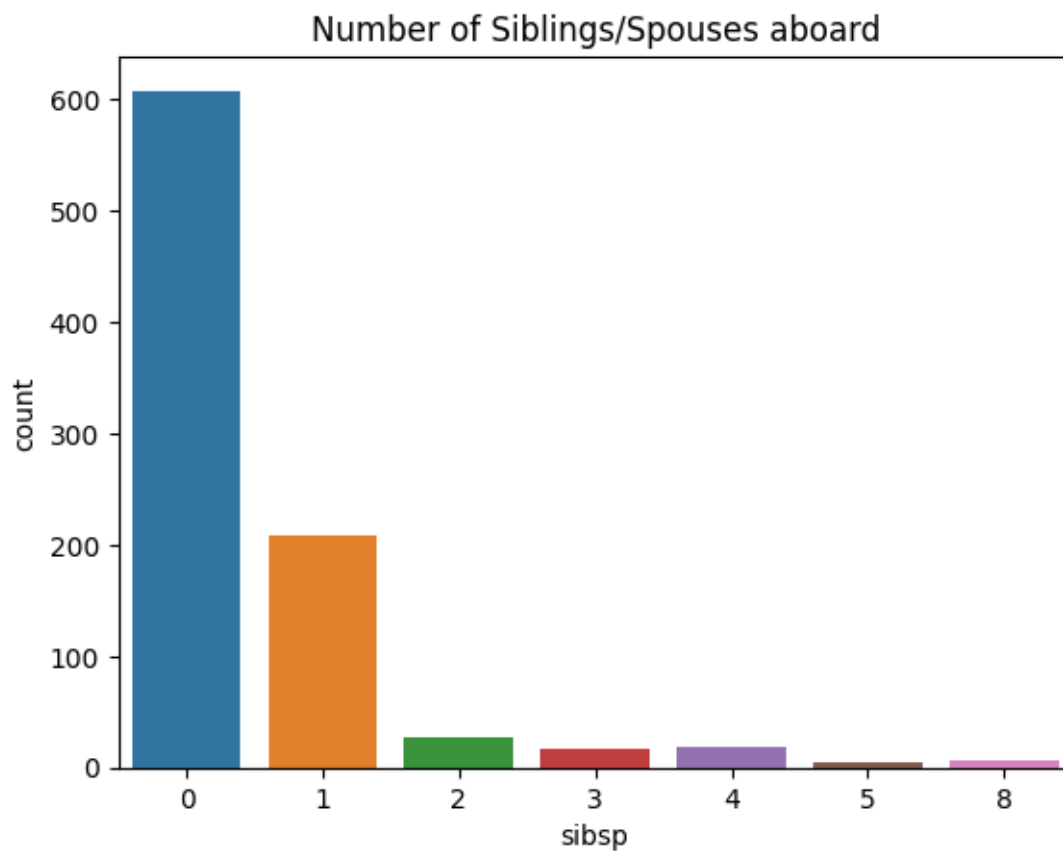
```
[216]: # Histogram for 'Fare'
df['fare'].plot(kind='hist', bins=30, color='lightgreen')
plt.title('Fare Distribution')
plt.show()
```




```
[178]: # Box plot for 'Fare'
sns.boxplot(x='survived', y='fare', data=df)
plt.title('Fare vs Survival')
plt.show()
```




```
[179]: # Count plot for 'SibSp'  
sns.countplot(x='sibsp', data=df)  
plt.title('Number of Siblings/Spouses aboard')  
plt.show()
```




```
[180]: # Survival rate by 'SibSp'  
sns.barplot(x='sibsp', y='survived', data=df)  
plt.title('Survival Rate by Siblings/Spouses aboard')  
plt.show()
```




```
[181]: # Count plot for 'Parch'  
sns.countplot(x='parch', data=df)  
plt.title('Number of Parents/Children aboard')  
plt.show()
```



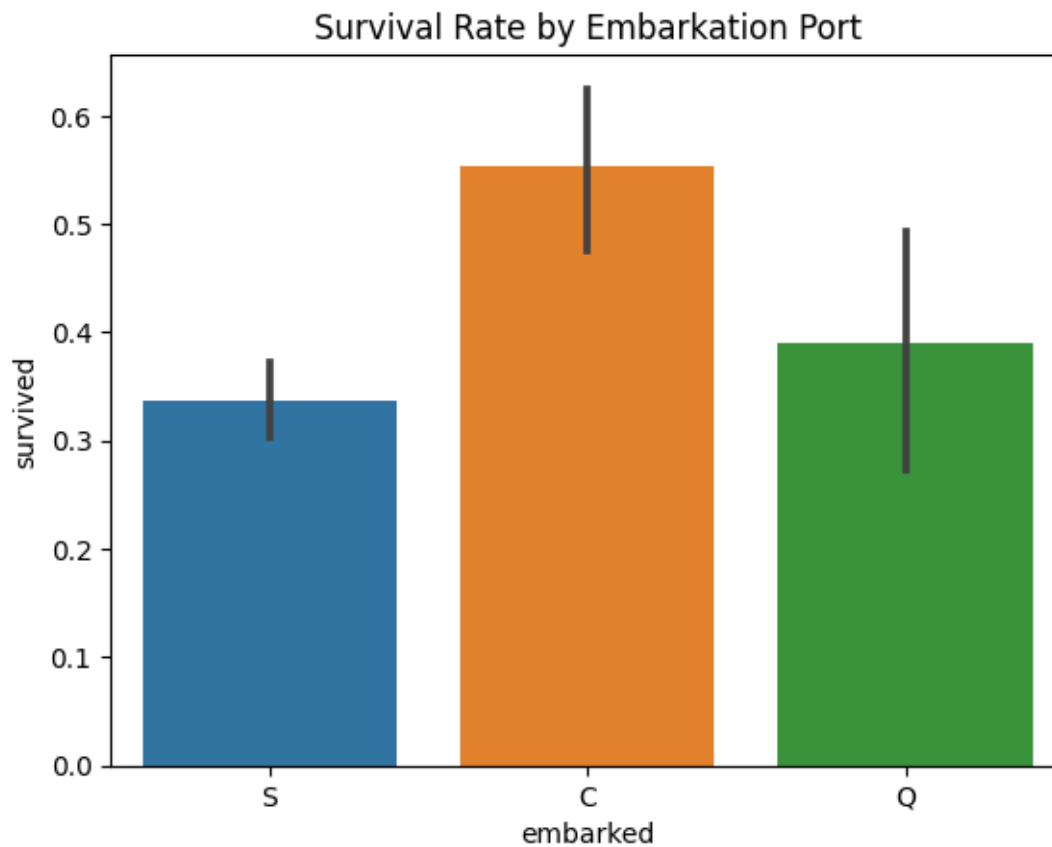
```
[182]: # Survival rate by 'Parch'
sns.barplot(x='parch', y='survived', data=df)
plt.title('Survival Rate by Parents/Children aboard')
plt.show()
```



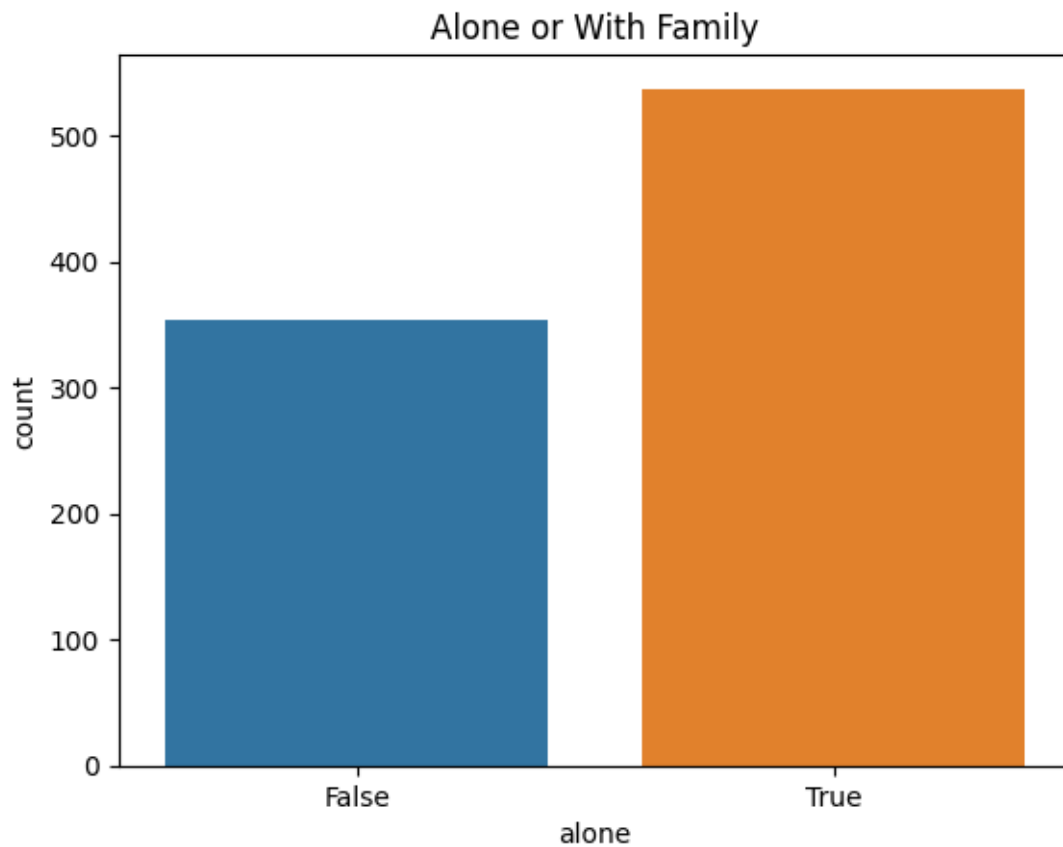
```
[183]: # Count plot for 'Embarked'  
sns.countplot(x='embarked', data=df)  
plt.title('Embarkation Port Distribution')  
plt.show()
```

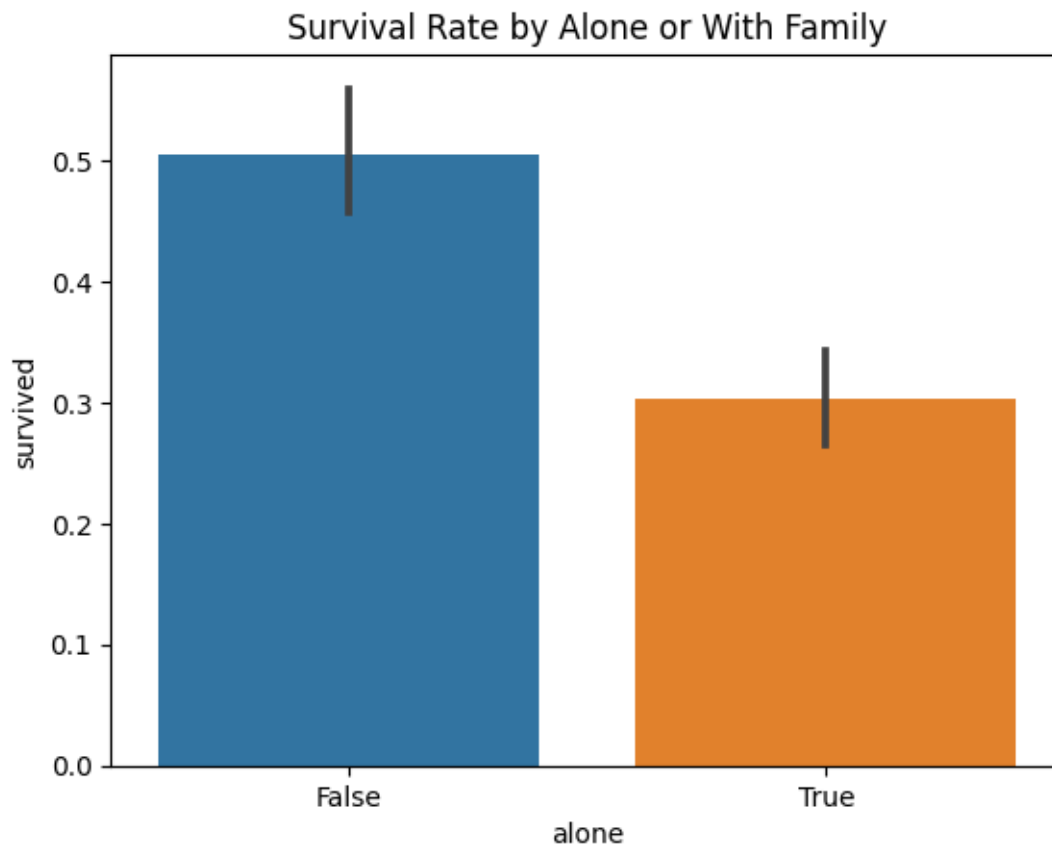
```
[184]: # Survival rate by 'Embarked'
sns.barplot(x='embarked', y='survived', data=df)
plt.title('Survival Rate by Embarkation Port')
plt.show()
```



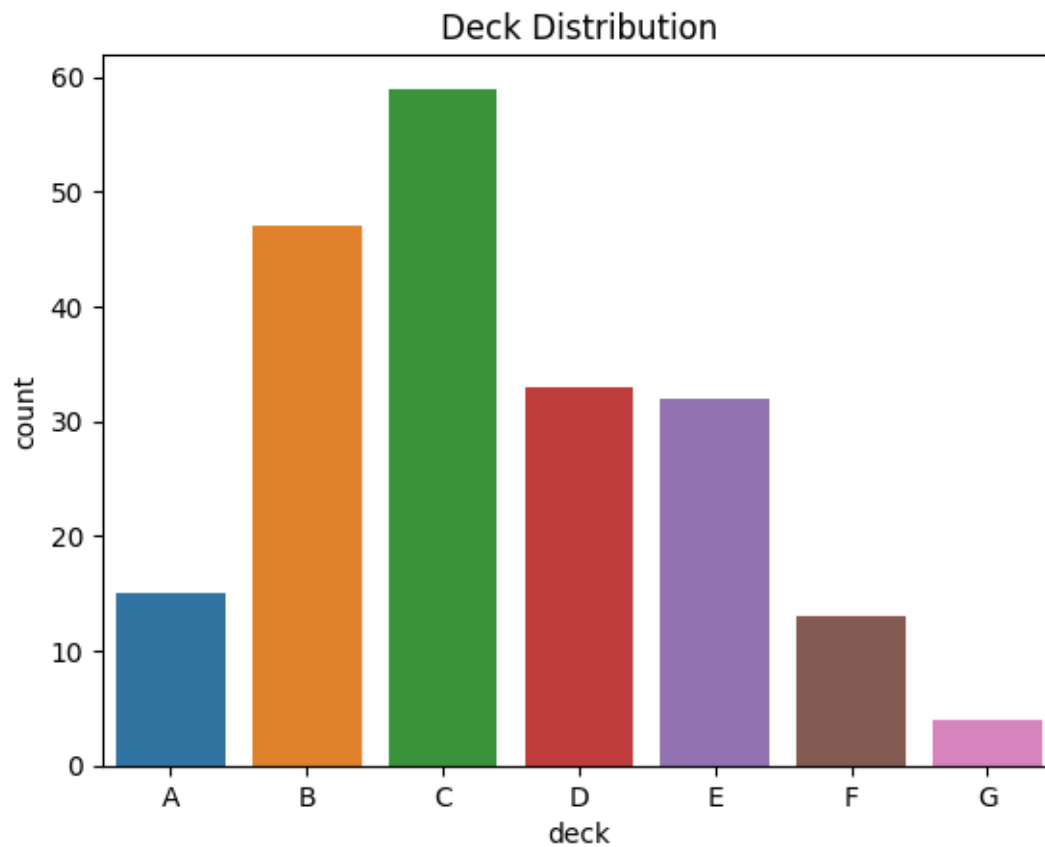
```
[185]: # Count plot for 'Alone'
sns.countplot(x='alone', data=df)
plt.title('Alone or With Family')
plt.show()
```




```
[186]: # Survival rate by 'Alone'
sns.barplot(x='alone', y='survived', data=df)
plt.title('Survival Rate by Alone or With Family')
plt.show()
```



```
[187]: # Count plot for 'Deck'
sns.countplot(x='deck', data=df)
plt.title('Deck Distribution')
plt.show()
```



```
[188]: # Survival rate by 'Deck'  
sns.barplot(x='deck', y='survived', data=df)  
plt.title('Survival Rate by Deck')  
plt.show()
```



```
[189]: # Fill missing 'age' based on median within 'pclass' and 'sex' groups
df['age'] = df.groupby(['pclass', 'sex'])['age'].transform(lambda x: x.fillna(x.
    ↪median()))
```

```
[190]: # Add 'Unknown' to the list of categories in 'deck' column
df['deck'] = df['deck'].cat.add_categories('Unknown')

# Fill missing 'deck' values with 'Unknown'
df['deck'].fillna('Unknown', inplace=True)
```

```
[191]: # Fill missing 'embarked' with mode
df['embarked'].fillna(df['embarked'].mode()[0], inplace=True)
```

```
[192]: # Fill missing 'embark_town' with mode
df['embark_town'].fillna(df['embark_town'].mode()[0], inplace=True)
```

```
[193]: df.isnull().sum()
```

```
[193]: survived      0
      pclass        0
      sex           0
      age           0
      sibsp         0
      parch         0
      fare          0
      embarked      0
      class         0
      who           0
      adult_male    0
      deck          0
      embark_town   0
      alive         0
      alone         0
      dtype: int64
```

```
[194]: # Encoding categorical variables
df['sex'] = df['sex'].map({'male': 0, 'female': 1})
df = pd.get_dummies(df, columns=['embarked', 'class', 'who', 'deck', 'embark_town'], drop_first=True)
```

```
[195]: # Create a 'family_size' feature
df['family_size'] = df['sibsp'] + df['parch']
```

```
[196]: # Drop irrelevant columns
df.drop(['alive', 'adult_male', 'alone'], axis=1, inplace=True)
```

```
[197]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   survived              891 non-null   int64
1   pclass                891 non-null   int64
2   sex                   891 non-null   int64
3   age                   891 non-null   float64
4   sibsp                 891 non-null   int64
5   parch                 891 non-null   int64
6   fare                  891 non-null   float64
7   embarked_Q            891 non-null   bool
8   embarked_S            891 non-null   bool
9   class_Second          891 non-null   bool
10  class_Third            891 non-null   bool
11  who_man                891 non-null   bool
```

```

12  who_woman                891 non-null    bool
13  deck_B                   891 non-null    bool
14  deck_C                   891 non-null    bool
15  deck_D                   891 non-null    bool
16  deck_E                   891 non-null    bool
17  deck_F                   891 non-null    bool
18  deck_G                   891 non-null    bool
19  deck_Unknown             891 non-null    bool
20  embark_town_Queenstown   891 non-null    bool
21  embark_town_Southampton  891 non-null    bool
22  family_size              891 non-null    int64
dtypes: bool(15), float64(2), int64(6)
memory usage: 68.9 KB

```

```

[198]: from sklearn.model_selection import train_test_split

X = df.drop('survived', axis=1)
y = df['survived']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

```

```

[199]: from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC

# Logistic Regression model
logreg = LogisticRegression()
logreg.fit(X_train, y_train)

# Random Forest model
rf = RandomForestClassifier()
rf.fit(X_train, y_train)

```

```

[199]: RandomForestClassifier()

```

```

[200]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

# Function to plot the confusion matrix for Titanic dataset
def plot_confusion_matrix(y_true, y_pred, model_name):
    cm = confusion_matrix(y_true, y_pred)

    plt.figure(figsize=(6, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,

```



```

        xticklabels=['Died', 'Survived'], yticklabels=['Died', 'Survived'])
    plt.title(f'{model_name} Confusion Matrix', fontsize=16)
    plt.xlabel('Predicted Label', fontsize=12)
    plt.ylabel('True Label', fontsize=12)
    plt.show()

```

```

[201]: from sklearn.metrics import accuracy_score, classification_report, \
        confusion_matrix

# Logistic Regression Evaluation
y_pred_logreg = logreg.predict(X_test)
print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred_logreg))
print(classification_report(y_test, y_pred_logreg))

```


Logistic Regression Accuracy: 0.8156424581005587

	precision	recall	f1-score	support
0	0.83	0.86	0.85	105
1	0.79	0.76	0.77	74
accuracy			0.82	179
macro avg	0.81	0.81	0.81	179
weighted avg	0.81	0.82	0.82	179

```

[202]: plot_confusion_matrix(y_test, y_pred_logreg, "Logistic Regression")

```




```
[203]: # Random Forest Evaluation
y_pred_rf = rf.predict(X_test)
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
print(classification_report(y_test, y_pred_rf))
```

```
Random Forest Accuracy: 0.7821229050279329
      precision    recall  f1-score   support

     0       0.82      0.81      0.81       105
     1       0.73      0.74      0.74        74

 accuracy          0.78          179
 macro avg         0.78          179
weighted avg         0.78          179
```

```
[204]: plot_confusion_matrix(y_test, y_pred_rf, "Random Forest Regression")
```



```
[206]: from sklearn.model_selection import GridSearchCV

# Example for Random Forest
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10]
}

# Perform GridSearchCV with Random Forest
grid_search_rf = GridSearchCV(estimator=rf, param_grid=param_grid, cv=3)
grid_search_rf.fit(X_train, y_train)

# Print the best parameters
print("Best parameters for Random Forest:", grid_search_rf.best_params_)
```

Best parameters for Random Forest: {'max_depth': 20, 'min_samples_split': 5, 'n_estimators': 200}

```
[211]: # Use the best parameters to create a new Random Forest model
best_rf = RandomForestClassifier(
    max_depth=grid_search_rf.best_params_['max_depth'],
    min_samples_split=grid_search_rf.best_params_['min_samples_split'],
```

```

    n_estimators=grid_search_rf.best_params_['n_estimators']
)

# Fit the model on the training data
best_rf.fit(X_train, y_train)

# Evaluate on the test data
y_pred = best_rf.predict(X_test)

# Accuracy
accuracy = best_rf.score(X_test, y_test)
print(f'Accuracy of the optimized Random Forest: {accuracy:.4f}')

```

Accuracy of the optimized Random Forest: 0.8101

```
[212]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.81	0.88	0.84	105
1	0.80	0.72	0.76	74
accuracy			0.81	179
macro avg	0.81	0.80	0.80	179
weighted avg	0.81	0.81	0.81	179

```
[209]: plot_confusion_matrix(y_test, y_pred, "Random Forest Optimized")
```

Random Forest Optimized Confusion Matrix

