

uber-price-prediction

September 26, 2024

```
[121]: # This Python 3 environment comes with many helpful analytics libraries
        ↳ installed
        # It is defined by the kaggle/python Docker image: https://github.com/kaggle/
        ↳ docker-python
        # For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list
↳ all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that
↳ gets preserved as output when you create a version using "Save & Run All"
# You can also write temporary files to /kaggle/temp/, but they won't be saved
↳ outside of the current session
```

/kaggle/input/uber-fares-dataset/uber.csv

```
[122]: import pandas as pd
import numpy as np

import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
warnings.filterwarnings("ignore", category=UserWarning)

df = pd.read_csv("/kaggle/input/uber-fares-dataset/uber.csv")
```

```
[123]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
```

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	200000 non-null	int64
1	key	200000 non-null	object
2	fare_amount	200000 non-null	float64
3	pickup_datetime	200000 non-null	object
4	pickup_longitude	200000 non-null	float64
5	pickup_latitude	200000 non-null	float64
6	dropoff_longitude	199999 non-null	float64
7	dropoff_latitude	199999 non-null	float64
8	passenger_count	200000 non-null	int64

dtypes: float64(5), int64(2), object(2)

memory usage: 13.7+ MB

```
[124]: # Dropping the missing values
df = df.dropna(subset=['dropoff_longitude', 'dropoff_latitude'])
```

```
[125]: # Convert pickup_datetime to datetime
df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])
```

```
[126]: # Extract day of the week, hour of the day, and month
df['day_of_week'] = df['pickup_datetime'].dt.dayofweek
df['hour_of_day'] = df['pickup_datetime'].dt.hour
df['month'] = df['pickup_datetime'].dt.month
```

```
[127]: # Drop the 'Unnamed: 0' and 'key' columns
df = df.drop(['Unnamed: 0', 'key'], axis=1)
```

```
[128]: import numpy as np

# Function to calculate haversine distance
def haversine(lon1, lat1, lon2, lat2):
    R = 6371 # Radius of Earth in kilometers
    lon1, lat1, lon2, lat2 = map(np.radians, [lon1, lat1, lon2, lat2])
    dlon = lon2 - lon1
    dlat = lat2 - lat1
    a = np.sin(dlat/2)**2 + np.cos(lat1) * np.cos(lat2) * np.sin(dlon/2)**2
    c = 2 * np.arctan2(np.sqrt(a), np.sqrt(1 - a))
    distance = R * c
    return distance

# Apply haversine distance calculation
df['distance_km'] = haversine(df['pickup_longitude'], df['pickup_latitude'],
                              df['dropoff_longitude'], df['dropoff_latitude'])
```

```
[129]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 199999 entries, 0 to 199999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fare_amount           199999 non-null  float64
1   pickup_datetime       199999 non-null  datetime64[ns, UTC]
2   pickup_longitude      199999 non-null  float64
3   pickup_latitude       199999 non-null  float64
4   dropoff_longitude     199999 non-null  float64
5   dropoff_latitude      199999 non-null  float64
6   passenger_count       199999 non-null  int64
7   day_of_week           199999 non-null  int32
8   hour_of_day           199999 non-null  int32
9   month                 199999 non-null  int32
10  distance_km           199999 non-null  float64
dtypes: datetime64[ns, UTC](1), float64(6), int32(3), int64(1)
memory usage: 16.0 MB
```

```
[131]: import numpy as np

def capOutliers(df, columns):
    df_capped = df.copy() # Create a copy of the original DataFrame

    for column in columns:
        lower_bound = df_capped[column].quantile(0.01) # Calculate the 1st
        ↪percentile
        upper_bound = df_capped[column].quantile(0.99) # Calculate the 99th
        ↪percentile

        # Cap the values in the column
        df_capped[column] = np.where(df_capped[column] < lower_bound,
        ↪lower_bound, df_capped[column])
        df_capped[column] = np.where(df_capped[column] > upper_bound,
        ↪upper_bound, df_capped[column])

    return df_capped

# Apply on fare_amount, passenger_count, distance_km
columns_to_cap = ["fare_amount", "passenger_count", "distance_km"]
df_cleaned = capOutliers(df, columns_to_cap)

print(f"Original data shape: {df.shape}")
print(f"Capped data shape: {df_cleaned.shape}")
```

```
Original data shape: (199999, 11)
Capped data shape: (199999, 11)
```

```
[132]: df_cleaned.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 199999 entries, 0 to 199999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fare_amount            199999 non-null  float64
1   pickup_datetime        199999 non-null  datetime64[ns, UTC]
2   pickup_longitude        199999 non-null  float64
3   pickup_latitude        199999 non-null  float64
4   dropoff_longitude       199999 non-null  float64
5   dropoff_latitude       199999 non-null  float64
6   passenger_count        199999 non-null  float64
7   day_of_week            199999 non-null  int32
8   hour_of_day            199999 non-null  int32
9   month                  199999 non-null  int32
10  distance_km            199999 non-null  float64
dtypes: datetime64[ns, UTC](1), float64(7), int32(3)
memory usage: 16.0 MB
```

```
[133]: import seaborn as sns
import matplotlib.pyplot as plt

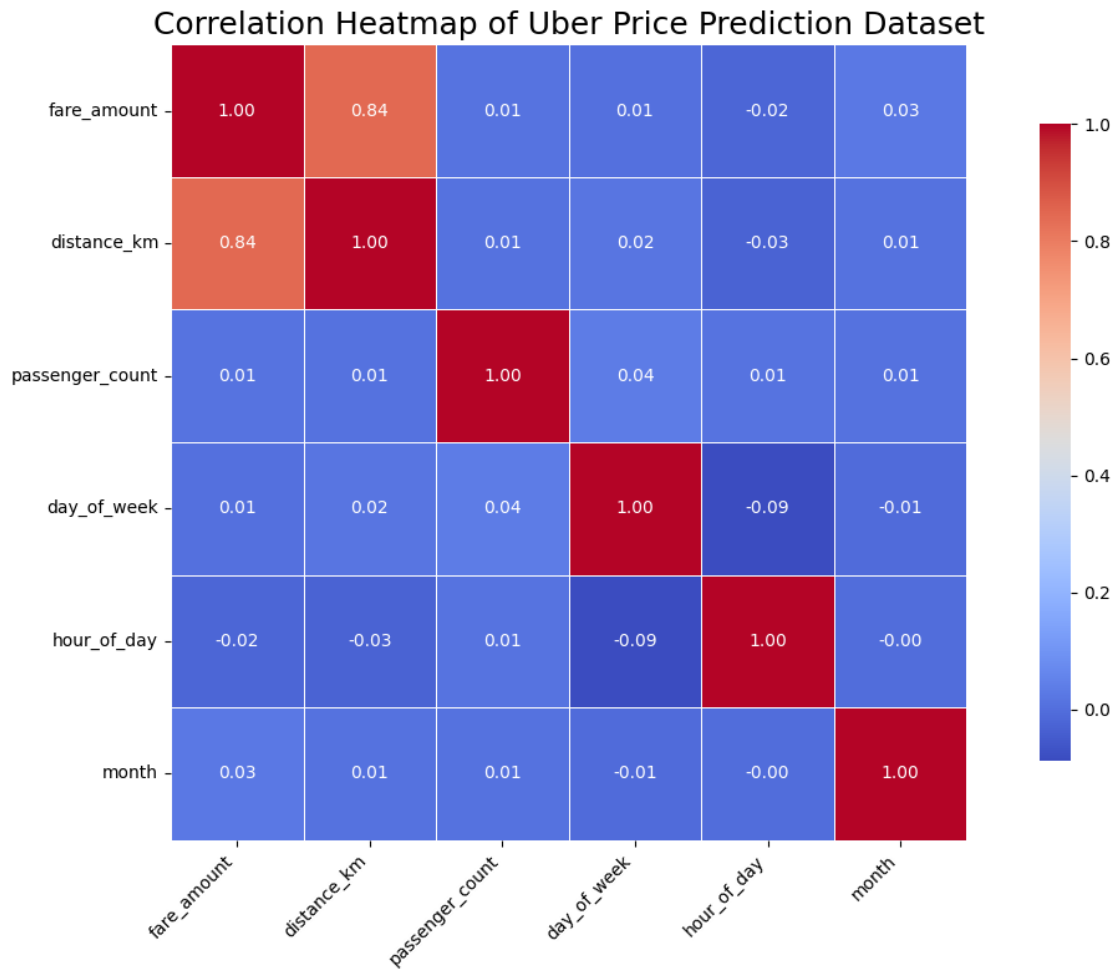
df_test = df_cleaned[["fare_amount", 'distance_km', 'passenger_count',
    ↪ 'day_of_week', 'hour_of_day', 'month']]
# Calculate the correlation matrix
correlation_matrix = df_test.corr()

# Set up the matplotlib figure
plt.figure(figsize=(12, 8))

# Create a heatmap
sns.heatmap(correlation_matrix,
            annot=True,          # Show correlation coefficients
            fmt=".2f",          # Format for the annotations
            cmap="coolwarm",    # Color palette
            square=True,        # Make cells square-shaped
            linewidths=0.5,     # Lines between cells
            cbar_kws={"shrink": .8}, # Color bar shrink
            mask=None)          # Masking for upper triangle if desired

# Title and labels
plt.title('Correlation Heatmap of Uber Price Prediction Dataset', fontsize=18)
plt.xticks(rotation=45, ha='right')
plt.yticks(rotation=0)
plt.tight_layout() # Adjust layout to make room for labels
```

```
# Show the plot
plt.show()
```



```
[134]: from sklearn.model_selection import train_test_split

# Define your features and target variable
X = df_cleaned[['distance_km', 'passenger_count', 'day_of_week', 'hour_of_day',
↪ 'month']]
y = df_cleaned['fare_amount']

# Split data into training and testing sets (80/20 split)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪ random_state=42)
```

```
[135]: from sklearn.linear_model import LinearRegression
```

```
# Initialize and train the model  
lin_reg = LinearRegression()  
lin_reg.fit(X_train, y_train)  
  
# Predict on the test set  
y_pred = lin_reg.predict(X_test)
```

```
[136]: from sklearn.metrics import mean_squared_error, r2_score
```

```
# Evaluate Linear Regression  
rmse = np.sqrt(mean_squared_error(y_test, y_pred))  
r2 = r2_score(y_test, y_pred)  
  
print(f'Linear Regression RMSE: {rmse}')
```

```
print(f'Linear Regression R2: {r2}')
```

Linear Regression RMSE: 5.00939230453408

Linear Regression R²: 0.6945436247499934

```
[137]: from sklearn.ensemble import RandomForestRegressor
```

```
# Initialize and train the model  
rf_reg = RandomForestRegressor(n_estimators=100, random_state=42)  
rf_reg.fit(X_train, y_train)  
  
# Predict on the test set  
y_pred_rf = rf_reg.predict(X_test)
```

```
[138]: # Evaluate Random Forest  
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))  
r2_rf = r2_score(y_test, y_pred_rf)
```

```
print(f'Random Forest RMSE: {rmse_rf}')
```

```
print(f'Random Forest R2: {r2_rf}')
```

Random Forest RMSE: 4.89074794162698

Random Forest R²: 0.7088413701984508